



IMC-4302C – Statistical learning

1. Linear regression and stochastic gradient descent

Kevin Zagalo

[<kevin.zagalo@inria.fr>](mailto:kevin.zagalo@inria.fr)

January 13, 2020

Use a **Jupyter** notebook and send it to [the depository](#) before the next session.

Any notebook sent after that won't be considered. Thank you !

Exercise 1 *Gradient descent*

We want to minimize a function $F : \mathbf{R}^d \rightarrow \mathbf{R}$, differential on \mathbf{R}^d . Let ∇F be its gradient. The gradient descent corresponds to the iterative algorithm :

- **Initialization** $x_0 \in \mathbf{R}^d$
- **Iteration** $x_{k+1} = x_k - \alpha \nabla F(x_k)$

The iteration is repeated until a stopping criterion is reached.

Question 1

Suggest some criterion to stop the algorithm.

Question 2

Implement this algorithm in Python. We will define a function which takes as input the function F and its gradient ∇F and others.

Question 3

What happens if F is not convex ?

Question 4

Discuss the influence of the x_0 point, in the non convex case, then in the convex case.

■

Exercise 2 *Least squared method*

A linear regression with n inputs $x_1, \dots, x_n \in \mathbf{R}^d$ and an output defined by $d + 1$ constants : the weights parameters w_1, \dots, w_d and a bias b . We define the regression function $h_w(x) = \langle x, w \rangle - b$.

We define the local and global error, respectively :

$$e(x; w) = \frac{1}{2}(y_x - h_w(x))^2 ; E(x_1, \dots, x_n; w) = \frac{1}{n} \sum_{i=1}^n e(x_i; w) \quad (1)$$

To estimate the final weight parameters β , we want to minimize the global error.

We want to implement the algorithm :

— **Input** $X = (x_1, \dots, x_n)$ and the associated responses (y_1, \dots, y_n) , and $\epsilon > 0$ a parameter.

— $t = 0$

— $w(t) = \vec{0}$

— **Repeat**

— **Compute** $L(w) = E(x_1, \dots, x_n; w)$

— **Update** $w(t+1) = w(t) - \epsilon \nabla L(w(t))$

Until all data is explored and *convergence* of the weights sequence to β .

Question 1

We treated the bias b as a parameter. Show that we can consider it as a weight parameter. How to adapt the problem of regression ?

Question 2

Write a function taking the weights parameters, an observation and its associated response that updates the weight parameters.

Question 3

Implement the algorithm.

Question 4

Plot the global error over the iterations using `matplotlib`.

Question 5

Modify the algorithm by implementing SGD and plot the global error over the iterations.

Question 6

Compare execution times of GD and SGD with the library `time`.



Exercise 3

Download [the `house.csv` file](#) containing 3 columns that represent the area, the number of rooms and the price of 600 houses (one per row). Comment every result.

Part 1 : Linear regression with 1 feature (house area)

In this first part, we will train a linear model for house price prediction using only one feature the house area. We will start by implementing a cost function and the gradient of this cost function. Then, we will implement the gradient descent algorithm that minimizes this cost function and determine the linear model parameter θ in the equation $h_{\theta}(x) = \theta_1 x$.

Question 1

Open this file with a file editor to understand more the data. Load the data in a `house_data` variable and check its size.

Hint: You could use [loadtxt](#) function from the `numpy` library.

Question 2

Extract m – the size of the sample and n – the number of features. Extract the house area and price columns respectively in X and y arrays to visualize them.

Hint: The shape of X and y arrays should be $(m, 1)$ for the following questions and not $(m,)$. You could use `newaxis` numpy object to add a new axis of length one.

Question 3

Implement the `cost_func` function that evaluate and return the previous equation of the mean squared error (1).

Question 4

Implement the `grad_cost_func` function that evaluates the gradient of the cost function at the point theta considering the j^{th} component as given on the previous equation.

Question 5

Implement the `grad_descent` algorithm that updates, iteratively, the parameter vector θ according to the previous equation.

Call the `grad_descent` function to compute θ_{opt} . You could use `max_iteration` of 1000 and α equal to 0.0001.

Use the calculated θ_{opt} to estimate the price of a house with $330m^2$ area.

Part 2 : Linear regression with 2 features (house area + bias term)

In this part, we will train a linear model for house price prediction using the house area and the bias term that represents the constant term (y-intercept) in the linear model equation $h_{\theta}(x) = \theta_1 x + \theta_0$.

Question 6

- Build the matrix X with shape $(m, 2)$ that represents 2 features: a column of ones that represents the bias term and a column of house area. (Use `numpy`'s `concatenate` function)
- Change the value of n to be equal to the number of features (number of columns of matrix X equal to 2 in this example).
- Make all needed modification in `cost_func`, `grad_cost_func` and `grad_descent` functions if your implementation was not generalizable for any number of features n .

Question 7

Use the calculated θ_{opt} of the new model to estimate the price of a house with $330m^2$ area.

Question 8

Try to normalize on the house area feature to enhance the convergence of the model. You should modify the X matrix in the previous code block and re-execute the code. What do you notice?

Part 3 : Linear regression with 3 features (house area + number of rooms + bias term)

In this part, we will train a linear model for house price prediction using the house area, number of rooms and the bias term that represents the constant term in the linear model equation $h_{\theta}(x) = \theta_2x_2 + \theta_1x_1 + \theta_0$.

Question 9

- Build the matrix X with shape $(m, 3)$ that represents 3 features: a column of ones that represents the bias term, a column of house area and a column of number of rooms.
- Change the value of n to be equal to the number of features (number of columns of matrix X equal to 3 in this example).
- Use the calculated θ_{opt} of the new model to estimate the price of a house with $330m^2$ area and 5 rooms. Compared to the previous model which predict better house prices ?

Bonus: You could also try to add other feature columns to the matrix X like area^2 or $\sqrt{\text{area}}$... and see the effect on the model and the error.

