



## IMC-4302C – Statistical learning

### 1. Linear regression and stochastic gradient descent

Kevin Zagalo

[<kevin.zagalo@inria.fr>](mailto:kevin.zagalo@inria.fr)

2019 – 2020

Use a Jupyter notebook and send it to [the depository](#) before the next session.

Any notebook sent after that won't be considered. Thank you !

### Exercise 1 *Gradient descent*

We want to minimize a function  $F : \mathbf{R}^d \rightarrow \mathbf{R}$ , differential on  $\mathbf{R}^d$ . Let  $\nabla F$  be its gradient. The gradient descent corresponds to the iterative algorithm :

— **Initialization**  $x_0 \in \mathbf{R}^d$

— **Iteration**  $x_{k+1} = x_k - \alpha \nabla F(x_k)$

The iteration is repeated until a stopping criterion is reached.

#### Question 1

Suggest some criterion to stop the algorithm.

#### Question 2

Implement this algorithm in Python. We will define a function which takes as input the function  $F$  and its gradient  $\nabla F$  and others.

### Question 3

What happens if  $F$  is not convex ?

### Question 4

Discuss the influence of the  $x_0$  point, in the non convex case, then in the convex case.

■

## Exercise 2 *Least squared method*

A linear regression with  $n$  inputs  $x_1, \dots, x_n \in \mathbf{R}^d$  and an output defined by  $d + 1$  constants : the weights parameters  $w_1, \dots, w_d$  and a bias  $b$ . We define the regression function  $h_w(x) = \langle x, w \rangle - b$ .

We define the local and global error, respectively :

$$e(x; w) = \frac{1}{2}(y_x - h_w(x))^2 ; E(x_1, \dots, x_n; w) = \frac{1}{n} \sum_{i=1}^n e(x_i; w) \quad (1)$$

To estimate the final weight parameters  $w^*$ , we want to minimize the global error.

We want to implement the algorithm :

— **Input**  $X = (x_1, \dots, x_n)$  and the associated responses  $(y_1, \dots, y_n)$ , and  $\epsilon > 0$  a parameter.

—  $t = 0$

—  $w(t) = \vec{0}$

— **Repeat**

— **Compute**  $L(w) = E(x_1, \dots, x_n; w)$

— **Update**  $w(t+1) = w(t) - \epsilon \nabla L(w(t))$

**Until** all data is explored and *convergence* of the weights sequence to  $\beta$ .

### Question 1

We treated the bias  $b$  as a parameter. Show that we can consider it as a weight parameter. How to adapt the problem of regression ?

### Question 2

Write a function taking the weights parameters, an observation and its associated response that updates the weight parameters.

### Question 3

Implement the algorithm.

### Question 4

Plot the global error over the iterations using `matplotlib`.

### Question 5

Modify the algorithm by implementing SGD and plot the global error over the iterations.

### Question 6

Compare execution times of GD and SGD with the library `time`.



## Exercise 3

Download [the `house.csv` file](#) containing 3 columns that represent the area, the number of rooms and the price of 600 houses (one per row). Comment every result.

### Question 1

Open this file with a file editor to understand more the data. Load the data and check its size.

**Hint:** You could use `loadtxt` function from the `numpy` library.

### Question 2

Extract the house area and price columns respectively in  $X$  and  $y$  lists. Scatter prices against areas.

### Question 3

The cost function we will use for this linear model training is the **Mean Squared Error** function defined by

$$L(w) = MSE(X \cdot w, y) = \frac{1}{2n} \sum_{i=1}^n (x_i \cdot w - y_i)^2$$

First of all, transform  $X$  and  $y$  to be respectively  $(1, n)$  and  $(n, 1)$ -numpy arrays.

#### Question 4

Implement the mean squared error cost function. Then implement its gradient :

$$\nabla L(w) = \partial_w MSE(X \cdot w, y) = \frac{1}{n} \sum_{i=1}^n (x_i \cdot w - y) x_i$$

#### Question 5

The update equation of the gradient descent algorithm is given by:

$$w^{(t+1)} = w^{(t)} - \alpha \nabla L(w^{(t)})$$

Where  $\alpha$  represents the step or the **learning rate**. Compute three steps of the gradient descent with different values of  $\alpha$ . For each  $\alpha$ , plot  $w \rightarrow L(w)$ , your initial point  $(w_0, L(w_0))$  and the three points associated to  $w_1$ ,  $w_2$  and  $w_3$  you've just computed. What the best  $\alpha$  according to you ?

#### Question 6

Implement the gradient descent algorithm, taking as input the gradient function, the learning rate, one or several stopping criterions and an initial parameter  $w_0$ .

#### Question 7

Compute the optimal parameter  $w^*$ . Compare on the same frame the actual prices and the ones predicted by the linear model against the areas.

#### Question 8

Introduce the bias term.

#### Question 9

Same as question 7.

#### Question 10

Add the number of rooms in  $X$ . Same as question 8.

\*Question 11 You could also add other features to the matrix  $X$  like  $area^{0.5}$ . Is it still linear ?

## Question 12

Create a LinearRegression class object as follows :

```
class LinearRegression:

    def __init__(self, alpha=1.0, tol=1e-3, max_iter=None, normalize=False):
        """
        Initializes the hyper parameters of the model
        alpha : Regularization strength; must be a positive float.
        max_iter : Number of steps of gradient descent.
        normalize : Whether to normalize features or not.
        solver : Gradient descent or Stochastic gradient descent
        """
        pass

    def fit(self, X, y):
        """
        Fit the linear model
        X : {ndarray, sparse matrix} of shape (n_samples, n_features)
        y : ndarray of shape (n_samples,) or (n_samples, n_targets)
        """
        pass

    def predict(self, X):
        """
        Predict using the linear model
        X : array_like or sparse matrix, shape (n_samples, n_features)
        """
        pass
```

## Question 13

How much would cost a  $330m^2$  flat with 5 rooms ?

